

Beyond Detection Parity: Localization Parity in Pedestrian Detection

Anwar Gashaw Yimam, Abel Adissu, Dawit Kindea, and Beakal Gizachew

Addis Ababa University, Addis Ababa, Ethiopia

`anwar.gashaw-ug@aau.edu.et`, `abel.adissu-ug@aau.edu.et`,
`dawit.kindea-ug@aau.edu.et`, `beakal.gizachew@aau.edu.et`

Abstract. Pedestrian detection is a safety-critical component of autonomous driving and intelligent surveillance, where equitable performance across demographic groups is essential for trustworthy deployment. Fairness research in this domain has evaluated performance almost exclusively through detection recall, that is, whether a person is found at all, leaving the precision of the predicted bounding box unexamined. Because existing audits reduce localization to a binary hit-or-miss threshold, systematic disparities in how tightly detectors localize different groups remain invisible, even though loose boxes directly compromise downstream distance and trajectory estimates. This paper introduces *Localization Parity*, a fairness criterion that requires equal expected Intersection over Union across demographic groups conditional on detection. We present the first confounder-controlled audit spanning four driving datasets (BDD100K, CityPersons, EuroCity-Day, and EuroCity-Night), three detector families (YOLOX, RetinaNet, and Cascade R-CNN), and two demographic axes (gender and age), together with a demographic-weighted bounding box regression intervention. Female and child pedestrians receive significantly looser boxes in all twelve dataset and detector combinations, with mean gaps of 1.40 and 3.27 percentage points. These gaps survive scale stratification in 35 of 36 and 24 of 30 cells, respectively. The proposed intervention closes the training-set gap by 46.8%, yet fails to generalize to validation data. Localization Parity therefore exposes a previously unmeasured and safety-relevant axis of algorithmic unfairness, and the resulting fairness generalization gap poses a concrete open problem for equitable perception systems.

Keywords: Localization Parity · Algorithmic Fairness · Pedestrian Detection · Bounding Box Regression · Demographic Bias · Object Detection

1 Introduction

Pedestrian detection is central to modern computer vision pipelines for self-driving cars, surveillance, and traffic infrastructure. The cost of error is not shared evenly across demographic groups. Prior work has shown that *detection*

recall is unequal: Wilson et al. [14] report lower average precision for darker-skinned pedestrians, while Brandao [2] and Li et al. [9] find substantially higher miss rates for children than for adults across many detectors.

What these audits leave unmeasured is *localization quality*. Once a detector draws a box, the tightness of that box shapes every estimate that follows: position, distance to the ego vehicle, predicted trajectory, and braking urgency. Two detections can both pass the conventional $\text{IoU} \geq 0.5$ hit threshold and still differ sharply in geometry. A box at IoU 0.51 and a box at IoU 0.91 look identical under recall-based fairness metrics, yet they feed very different signals into downstream modules. If one demographic group systematically receives looser boxes, that group systematically receives worse downstream estimates, and no current fairness benchmark would notice.

We formalize this overlooked criterion as **Localization Parity (LP)**: equality of expected Intersection over Union (IoU) across demographic groups, conditional on successful detection and after controlling for confounders such as object scale. We then run the first multi-dataset, multi-detector LP audit and test a training-time mitigation.

Contributions.

1. A formal definition of Localization Parity as a group fairness criterion on box regression quality, orthogonal to detection parity, with a clear link to existing fairness notions for classification and regression (Section 3).
2. The first confounder-controlled LP audit across four autonomous driving datasets, three detector families spanning single-stage and two-stage designs, and two demographic axes: *gender* (female vs. male) and *age* (child vs. adult). Race and skin tone are not audited in this study (Section 6.1).
3. Empirical evidence that female and child pedestrians receive consistently lower IoU boxes (mean gaps -1.40 pp and -3.27 pp), surviving ICON² scale stratification [13].
4. Demographic-Weighted Bounding Box Regression (DWBBR), which closes the training gap by 46.8% yet fails to transfer cleanly to held-out data. We document this *fairness generalization gap* as an open problem.

2 Related Work

Fairness in computer vision. Group fairness (demographic parity and equalized odds) [7], individual fairness, and calibration form the classical toolkit. Vision audits usually operationalize group fairness through equal detection rates or equal average precision. Wilson et al. [14] established predictive inequity across Fitzpatrick skin-tone groups on BDD100K using AP. Schumann et al. [12] critique binary third-party gender labels. Fair regression definitions [1] constrain disparities in continuous prediction error; LP applies the same idea to the IoU of matched detections (Section 3).

Pedestrian detection and demographic bias. BDD100K [15], CityPersons [16], and EuroCity Persons [3] are the dominant driving benchmarks. Li et al. [9] and Brandao [2] show elevated child miss rates across many detectors. Khoshkdahan et al. [8] study pose and occlusion as structural confounders of detection accuracy. All of these works evaluate hit-or-miss outcomes rather than continuous localization quality.

Confounder control and fair training. ICON² [13] stratifies by scale and occlusion so that demographic gaps cannot be attributed to size alone; we apply it to IoU rather than AP. Loss reweighting and Group DRO [11] are standard tools for fair training, but they have been used mainly for classification or recall rather than for the bounding box regression branch. Focal Loss [10] reweights hard classification examples; DWBBR instead reweights the CIoU regression term by demographic group performance.

3 Localization Parity

3.1 Formal Definition

Let $\mathcal{D} = \{(b_i, \hat{g}_i, g_i)\}$ be the set of successfully detected instances, where b_i is the predicted box, $\hat{g}_i$ the ground-truth box, and $g_i \in \{g_1, \dots, g_k\}$ the demographic group of instance i . Matching follows the COCO convention: a prediction counts as a detection only if $\text{IoU}(b_i, \hat{g}_i) \geq 0.5$. Localization Parity requires

$$\mathbb{E}[\text{IoU}(b_i, \hat{g}_i) \mid g_i = g_a] = \mathbb{E}[\text{IoU}(b_i, \hat{g}_i) \mid g_i = g_b] \quad (1)$$

for every pair of groups $a \neq b$, after controlling for instance scale. We report the signed gap

$$\Delta\text{IoU} = \mathbb{E}[\text{IoU} \mid g_{\text{priv}}] - \mathbb{E}[\text{IoU} \mid g_{\text{dis}}],$$

where g_{priv} denotes the privileged group (male, adult) and g_{dis} the disadvantaged group (female, child). A negative gap means the disadvantaged group receives worse boxes.

LP is deliberately conditional on detection. It equalizes box quality among people who were found, and is therefore orthogonal to detection parity (equal recall or miss rate). A detector can satisfy one while violating the other; both are needed for safety-relevant fairness.

3.2 Relation to Existing Fairness Criteria

LP sits at the intersection of group fairness and fair regression:

- Relative to demographic parity and equalized odds for classification [7, 5], LP replaces a binary outcome (detected or not) with a continuous quality score (IoU), evaluated only on the detected set. This is analogous to equalized opportunity applied to localization quality.

- Relative to fair regression [1], which seeks bounded disparity in expected loss across groups, LP uses the loss $\ell = 1 - \text{IoU}$ on matched boxes.
- Relative to AP-based audits [14], LP isolates the localization component that AP mixes with classification confidence and ranking.

LP does not claim individual fairness or calibration. It is a group-level criterion because the safety harm of interest, systematically worse downstream estimates for a demographic group, is itself group structured.

3.3 Why Small IoU Gaps Matter Downstream

A natural question is whether gaps of 1–3 pp IoU are practically consequential. Under a pinhole projection, relative depth error scales with relative box-height error. For a pedestrian at 30 m whose true box height is 80 px, a 2 pp mean IoU deficit (roughly 2–4 px of systematic height or vertical-center error) translates into about 0.7–1.5 m of depth bias. That is enough to alter time-to-collision estimates at urban speeds and to shift predicted crossing trajectories. Because trajectory and collision modules are typically nonlinear in box geometry, group-level mean gaps of this magnitude can compound across the planning stack even when every detection still clears the 0.5 IoU threshold. We therefore treat LP gaps in this range as safety-relevant, not merely statistically detectable. A full end-to-end propagation study remains future work (Section 8).

4 Experimental Setup

4.1 Datasets

We evaluate on four pedestrian benchmarks chosen for geographic, sensor, and lighting diversity:

- **BDD100K** [15]: large-scale US dashcam data with about 100,000 annotated pedestrian boxes and binary male/female labels.
- **CityPersons** [16]: about 19,000 boxes over 2,975 Cityscapes street-level images, with careful occlusion annotation.
- **EuroCity-Day** and **EuroCity-Night** [3]: subsets of EuroCity Persons (about 238,200 boxes over 47,300 images from 31 European cities), split by lighting.

Demographic labels for CityPersons and EuroCity use augmented annotations where available. Detections with missing or ambiguous labels are dropped. Label quality limitations are discussed in Section 7.

4.2 Detectors

We cover the three main detector design families with public pretrained weights and no fine-tuning in the audit phase:

- **YOLOX** [6]: single-stage, anchor-free;
- **RetinaNet** [10]: single-stage, anchor-based;

- **Cascade R-CNN** [4]: two-stage, with cascaded IoU refinement aimed at tighter localization.

The mitigation experiments use **YOLOv8n** (Ultralytics), the lightweight YOLO variant that serves as our trainable single-stage backbone.

4.3 Demographic Axes

The two protected attributes studied are:

1. **Gender:** female vs. male;
2. **Age:** child vs. adult.

Race and skin tone are not among the audited axes. An earlier pilot examined skin tone, but the labeled sample was too small and the categories too coarse for reliable statistical claims. We therefore exclude it here and mark it as future work (Section 7).

4.4 Audit Protocol and Confounder Control

Matching. Predictions are matched to ground truth by greedy, confidence-sorted IoU assignment. Matches with $\text{IoU} \geq 0.5$ enter the LP analysis; false positives are excluded.

Confounders. The primary confounder is **object scale** (box height), because female and child pedestrians tend to occupy smaller image regions, and small objects are harder to localize for any reason. Where annotations permit, we also consider **occlusion**. Following ICON² [13], we bin ground-truth instances within each dataset into short (≤ 33 rd height percentile), medium (34th–66th), and tall (≥ 67 th) strata, then recompute per-group mean IoU inside each bin. A gap that remains significant within strata cannot be attributed to scale alone.

Statistical testing. Within each cell and each stratum we compare group IoU distributions with the two-sided Mann–Whitney U test. IoU distributions are non-Gaussian and heavy-tailed, so a non-parametric rank test is appropriate. Significance is declared at $p < 0.05$. We report both raw and stratified gaps.

5 Demographic-Weighted Bounding Box Regression

DWBRR applies the Group DRO reweighting principle [11] to the CIoU regression branch. Standard YOLO training minimizes $L_{\text{class}} + \lambda L_{\text{reg}}$. DWBRR replaces the regression term with a group-weighted sum:

$$L_{\text{DWBRR}} = L_{\text{class}} + \lambda \sum_i w_{g_i} L_{\text{CIoU}}(b_i, \hat{g}_i), \quad (2)$$

where $w_g \propto 1/\mathbb{E}_{\text{train}}[\text{IoU} \mid g]$ and weights are renormalized each epoch so that $\sum_g w_g = 1$. No architecture change is required. We override the Ultralytics training callback to recompute weights from per-group training CIoU and scale the regression loss before backpropagation.




Fig. 1. Conceptual framework. Phase 1 audits Localization Parity across four datasets and three detectors under ICON² confounder control. Phase 2 and Phase 3 test whether DWBBR closes the gap and reveal a fairness generalization barrier.

Training protocol. Both DWBBR and an identically configured baseline (unweighted YOLOv8n) are trained on a 1,693-image BDD100K subset with reliable gender labels (64.8% male and 35.2% female instances), for 50 epochs, batch size 8, image size 416×416 , $\text{lr}_0 = 0.01$ with cosine decay, on a single NVIDIA RTX 3050. Validation uses a disjoint 423-image subset. We replicate with seeds {42, 123, 7}.

6 Results

6.1 Phase 1: Localization Parity Audit

Figure 1 summarizes the study design. Tables 1 and 2 and Figure 2 report the Phase 1 audit.

Gender. Female pedestrians receive lower mean IoU than male pedestrians in all 12 dataset and detector cells (mean gap -1.40 pp; 11 of 12 significant). The lone non-significant cell is CityPersons with Cascade R-CNN (-0.49 pp, $p = 0.11$). BDD100K shows the most consistent gaps (-2.16 , -2.02 , and -2.32 pp). EuroCity-Night is comparably severe.

Age. Children receive lower mean IoU than adults in all 12 cells (mean -3.27 pp; 11 of 12 significant), roughly double the gender gaps. EuroCity-Night is the harshest setting across all three detectors (-4.35 , -4.34 , and -4.44 pp).

Confounder control. Under ICON² scale stratification the gender gap survives in 35 of 36 cells (mean -1.43 pp, nearly identical to the unstratified -1.40 pp). The age gap survives in 24 of 30 within-scale cells. The disparities are therefore not reducible to women and children occupying smaller image regions.

Table 1. Phase 1 LP audit, gender axis. Gap = female minus male mean IoU (pp). †: $p < 0.05$, Mann–Whitney U . All 12 gaps are negative; 11 of 12 are significant.

Dataset	Detector	Gap (pp)
BDD100K	YOLOX	$-2.16^\dagger$
BDD100K	RetinaNet	$-2.02^\dagger$
BDD100K	Cascade R-CNN	$-2.32^\dagger$
CityPersons	YOLOX	$-0.98^\dagger$
CityPersons	RetinaNet	$-0.71^\dagger$
CityPersons	Cascade R-CNN	-0.49
EuroCity-Day	YOLOX	$-0.67^\dagger$
EuroCity-Day	RetinaNet	$-0.65^\dagger$
EuroCity-Day	Cascade R-CNN	$-0.98^\dagger$
EuroCity-Night	YOLOX	$-1.81^\dagger$
EuroCity-Night	RetinaNet	$-2.47^\dagger$
EuroCity-Night	Cascade R-CNN	$-1.59^\dagger$
Mean		-1.40

6.2 Phase 2 and Phase 3: DWBBR Intervention

Training dynamics. On the training set, DWBBR behaves as intended (Figure 3): the female to male CIoU gap narrows from -2.82 pp at epoch 1 to -1.50 pp at epoch 50 (46.8% relative reduction), while the female loss weight drifts toward parity as the per-group IoU ratio improves.

Validation. On held-out data the picture changes (Table 3, Figure 4). Under the primary seed, DWBBR female CIoU is 0.7715 against baseline 0.7790; the female to male gap widens from -1.52 pp (baseline) to -2.74 pp (DWBBR). A one-sided Mann–Whitney test finds no significant female improvement ($p = 0.71$). Across three seeds the mean validation gap moves from -1.78 ± 0.48 pp (baseline) to -1.36 ± 0.63 pp (DWBBR), a 23.6% mean reduction accompanied by higher variance and unstable best-checkpoint epochs ($\{15, 17, 21\}$ for DWBBR versus $\{34, 17, 24\}$ for baseline).

Fairness generalization gap. DWBBR reliably closes the training LP gap, yet the improvement does not transfer cleanly to validation: variance rises, checkpoint selection becomes unstable, and under the primary seed the validation gap worsens. This aligns with known limits of group reweighting under train and test demographic shift [11]. We report it as a finding rather than as a claim that all mitigations must fail, and we treat fairness-aware checkpoint selection and stronger DRO-style objectives as concrete next steps. Figure 5 shows qualitative Localization Parity violations on BDD100K.

Table 2. Phase 1 LP audit, age axis. Gap = child minus adult mean IoU (pp). †: $p < 0.05$, Mann–Whitney U . All 12 gaps are negative; 11 of 12 are significant.

Dataset	Detector	Gap (pp)
BDD100K	YOLOX	−1.92 [†]
BDD100K	RetinaNet	−1.64
BDD100K	Cascade R-CNN	−2.83 [†]
CityPersons	YOLOX	−2.91 [†]
CityPersons	RetinaNet	−3.11 [†]
CityPersons	Cascade R-CNN	−3.48 [†]
EuroCity-Day	YOLOX	−4.55 [†]
EuroCity-Day	RetinaNet	−2.49 [†]
EuroCity-Day	Cascade R-CNN	−3.17 [†]
EuroCity-Night	YOLOX	−4.35 [†]
EuroCity-Night	RetinaNet	−4.34 [†]
EuroCity-Night	Cascade R-CNN	−4.44 [†]
Mean		−3.27

Table 3. Validation CIoU by group on the BDD100K validation subset (423 images). Gap = weaker group minus reference group (pp). Multi-seed means use seeds {42, 123, 7}.

Group	Model	n	CIoU	95% CI	Gap (pp)
Male	Baseline	555	0.7942	—	—
Female	Baseline	376	0.7790	[0.768, 0.790]	−1.52
Male	DWBRR	555	0.7990	—	—
Female	DWBRR	374	0.7715	[0.760, 0.783]	−2.74
Mean female to male, baseline					−1.78 ± 0.48
Mean female to male, DWBRR					−1.36 ± 0.63

7 Ethical Considerations

Demographic labels in BDD100K are third-party visual inferences: binary gender, no self-identification, and undocumented annotation guidelines. These raise consent, accuracy, and inclusivity concerns [12]. LP is a proxy for box-quality fairness, not a guarantee of equal downstream risk. Nonlinear collision estimators can map equal mean IoU to unequal outcomes. Fairer localization can reduce collision risk and can also improve re-identification and surveillance. We publish because we judge the near-term safety benefit to outweigh that dual-use risk, while acknowledging that the judgment is contestable.

The study examines only gender and age. We do not report race or skin-tone LP results here. Intersectional analyses (for example, young female pedestrians) and broader attributes such as disability and body type remain open. Binary gender labels also exclude non-binary identities from the analysis.




Fig. 2. Phase 1 LP audit: gender gaps (a) and age gaps (b) across four datasets and three detectors. Asterisks mark $p < 0.05$.




Fig. 3. DWBBR training dynamics: LP gap convergence over 50 epochs on BDD100K.




Fig. 4. Validation-set CioU for DWBBR and the baseline, by demographic group.

8 Conclusion

Localization Parity is a previously unmeasured, safety-relevant fairness criterion for pedestrian detection. Across four datasets and three detector families, female



Fig. 5. Qualitative Localization Parity violations on BDD100K. Female pedestrians more often receive looser predicted boxes than male pedestrians in comparable scenes.

and child pedestrians receive systematically lower IoU boxes, and the gaps survive scale stratification. Demographic-weighted regression closes the training gap by 46.8% but does not generalize reliably. That result establishes the fairness generalization gap as an open research problem.

Limitations. The intervention is evaluated on BDD100K with YOLOv8n only. The validation child cohort is small ($n = 17/16$). Demographic labels are third-party and binary. We do not yet quantify end-to-end propagation of IoU gaps into planning errors.

Future work. Promising next steps include cross-dataset and cross-architecture DWBBR evaluation, fairness-aware checkpoint selection, Group DRO and over-sampling baselines, intersectional and skin-tone analyses, and closed-loop simulation of how 1–3 pp IoU disparities propagate into distance, time-to-collision, and braking decisions, including domains beyond autonomous driving such as surveillance and robotics.

Acknowledgements We thank the PanAfriCon AI 2026 Program Committee and reviewers for constructive feedback that improved the clarity of the experimental setup, the theoretical positioning of Localization Parity, and the discussion of practical safety implications.

References

1. Agarwal, A., Dudík, M., Wu, Z.S.: Fair regression: Quantitative definitions and reduction-based algorithms. Proceedings of the International Conference on Machine Learning (ICML) (2019)

2. Brandao, M.: Age and gender bias in pedestrian detection algorithms. arXiv preprint arXiv:1906.10490 (2019)
3. Braun, M., Krebs, S., Flohr, F.B., Gavrilu, D.M.: EuroCity Persons: A novel benchmark for person detection in traffic scenes. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **41**(8), 1844–1861 (2019)
4. Cai, Z., Vasconcelos, N.: Cascade R-CNN: High quality object detection and instance segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **43**(5), 1483–1498 (2021)
5. Calders, T., Verwer, S.: Three naive Bayes approaches for discrimination-free classification. In: *Data Mining and Knowledge Discovery* (2010)
6. Ge, Z., Liu, S., Wang, F., Li, Z., Sun, J.: YOLOX: Exceeding YOLO series in 2021. arXiv preprint arXiv:2107.08430 (2021)
7. Hardt, M., Price, E., Srebro, N.: Equality of opportunity in supervised learning. In: *Advances in Neural Information Processing Systems (NeurIPS)* (2016)
8. Khoshkdahan, M., Akbari, A., Akbari, A., Zhang, X.: Beyond overall accuracy: Pose- and occlusion-driven fairness analysis in pedestrian detection for autonomous driving. arXiv preprint (2025)
9. Li, X., Chen, Z., Zhang, J.M., Sarro, F., Zhang, Y., Liu, X.: Bias behind the wheel: Fairness testing of autonomous driving systems. *ACM Transactions on Software Engineering and Methodology* **34**(3), 1–24 (2025)
10. Lin, T.Y., Goyal, P., Girshick, R., He, K., Dollár, P.: Focal loss for dense object detection. In: *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*. pp. 2980–2988 (2017)
11. Sagawa, S., Koh, P.W., Hashimoto, T.B., Liang, P.: Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization. In: *International Conference on Learning Representations (ICLR)* (2020)
12. Schumann, C., Wang, S., Russakovsky, O., Sohl-Dickstein, J., Moritz, P., Forde, J.: A step toward more inclusive people annotations for fairness. In: *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* (2021)
13. Sudhakar, S., Prabhu, V., Russakovsky, O., Hoffman, J.: ICON²: Reliably benchmarking predictive inequity in object detection. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)* (2023)
14. Wilson, B., Hoffman, J., Morgenstern, J.: Predictive inequity in object detection. arXiv preprint arXiv:1902.11097 (2019)
15. Yu, F., Chen, H., Wang, X., Xian, W., Chen, Y., Liu, F., Madhavan, V., Darrell, T.: BDD100K: A diverse driving dataset for heterogeneous multitask learning. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. pp. 2636–2645 (2020)
16. Zhang, S., Benenson, R., Schiele, B.: CityPersons: A diverse dataset for pedestrian detection. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. pp. 3213–3221 (2017)